

Homework 3

BBIC: Bovee, Bichay, Iseneker, Coeman

2/8/2018

Problem 1

OLS function

First we simulate data for our LM

```
### simulate data

library(MASS)
set.seed(6886)
data = mvrnorm(
  n=100, mu=c(-2, 3),
  Sigma=matrix(c(8,3,3,2),nrow=2,ncol=2)
)
colnames(data) = c('y','x')
```

Here we create the function that calculates an OLS model manually. Comments within explain each step.

```
### create function

my_lm <- function(q, p){

  # load data

  x = as.matrix(cbind(1,p))
  y = as.matrix(q)

  # calculate b

  beta = solve( t(x) %*% x ) %*% t(x) %*% y

  # calculate vcv and se

  yhat <- x %*% beta # predicted values

  ssr <- sum((yhat - y)^2) # sum of squared residuals

  # set N = number of observations; k = number of variables (incl. intercept)

  N <- nrow(y)
  k <- ncol(x)

  sigma2 <- (ssr/(N-k)) # variance

  vcv <- (sigma2)*(solve( t(x) %*% x )) # vcv
```

```

se <- sqrt(diag(vcv))

# output results

model_summary <- matrix(c(beta, se), nrow=2, ncol=2,
                        dimnames=list( c("b0", "b1"),
                                       c("Coef", "S.E."))
                        )

return(model_summary)

}

my_lm(data[, "y"], data[, "x"])

##           Coef           S.E.
## b0 -6.390654 0.4681754
## b1  1.499807 0.1309778

```

Problem 2

“Create this” matrix, from slide 37

Use the rep function to fill the data in the data frame.

```

clever_df <- data.frame(Patient = c(1:16),
                        Gender = rep(c("Male", "Female"), each = 8),
                        Treatment1 = rep(c("Yes", "No"), each = 4),
                        Treatment2 = rep(c("Yes", "No"), each = 2),
                        Treatment3 = rep(c("Yes", "No"))
                        )

clever_df

```

	Patient	Gender	Treatment1	Treatment2	Treatment3
## 1	1	Male	Yes	Yes	Yes
## 2	2	Male	Yes	Yes	No
## 3	3	Male	Yes	No	Yes
## 4	4	Male	Yes	No	No
## 5	5	Male	No	Yes	Yes
## 6	6	Male	No	Yes	No
## 7	7	Male	No	No	Yes
## 8	8	Male	No	No	No
## 9	9	Female	Yes	Yes	Yes
## 10	10	Female	Yes	Yes	No
## 11	11	Female	Yes	No	Yes
## 12	12	Female	Yes	No	No
## 13	13	Female	No	Yes	Yes
## 14	14	Female	No	Yes	No
## 15	15	Female	No	No	Yes
## 16	16	Female	No	No	No

Problem 3

Polity matrix and summary stats

We first load the data, and then store the various statistics in respective vectors

```
load("/Users/zacharycoeman/Desktop/900/polity_dataframe.rda")

mean <- c(mean(polity$democ, na.rm=TRUE), mean(polity$autoc, na.rm=TRUE),
          mean(polity$polity2, na.rm=TRUE), mean(polity$xconst, na.rm=TRUE))

median <- c(median(polity$democ, na.rm=TRUE), median(polity$autoc, na.rm=TRUE),
            median(polity$polity2, na.rm=TRUE), median(polity$xconst, na.rm=TRUE))

stdev <- c(sd(polity$democ, na.rm=TRUE), sd(polity$autoc, na.rm=TRUE),
           sd(polity$polity2, na.rm=TRUE), sd(polity$xconst, na.rm=TRUE))

max <- c(max(polity$democ, na.rm=TRUE), max(polity$autoc, na.rm=TRUE),
         max(polity$polity2, na.rm=TRUE), max(polity$xconst, na.rm=TRUE))

min <- c(min(polity$democ, na.rm=TRUE), min(polity$autoc, na.rm=TRUE),
        min(polity$polity2, na.rm=TRUE), min(polity$xconst, na.rm=TRUE))

n <- c(NROW(na.omit(polity$democ)), NROW(na.omit(polity$autoc)),
      NROW(na.omit(polity$polity2)), NROW(na.omit(polity$xconst)))

nNA <- c(sum(is.na(polity$democ)), sum(is.na(polity$autoc)),
        sum(is.na(polity$polity2)), sum(is.na(polity$xconst)))
```

Next, we create a matrix from the 7 vectors

```
summary <- matrix(c(n, nNA, mean, median, max, min, stdev), nrow=4, ncol=7,
                  dimnames=list(names(polity[5:8]),
                                c("N", "nNA", "Mean", "Median", "Max", "Min", "StDev")))
)
```

And finally, we round the matrix to two decimal places, and transpose it so it matches the matrix in the slide

```
summary <- round(t(summary), 2)
```

```
summary
```

##	democ	autoc	polity2	xconst
## N	8122.00	8122.00	8385.00	8122.00
## nNA	357.00	357.00	94.00	357.00
## Mean	4.27	3.38	0.85	4.19
## Median	3.00	2.00	1.00	4.00
## Max	10.00	10.00	10.00	7.00
## Min	0.00	0.00	-10.00	1.00
## StDev	4.17	3.57	7.43	2.33

Problem 4

Merging the other variables

We first load in the two datasets

```
worldBank <- read.csv( file='/Users/zacharycoeman/Desktop/900/worldBank.csv', stringsAsFactors=FALSE
)

load("/Users/zacharycoeman/Desktop/900/polity_dataframe.rda")

colnames(worldBank)[4:7] = c( 'gdpGrowth', 'gdppcConstant', 'popGrowth', 'tradeOpen'
)
```

Next, we use the package countrycode to make a matchable country name variable, and then account for the time-series nature of the data by creating a country-year indicator

```
library(countrycode)
polity$cname <- countrycode(
  sourcevar=polity$country,
  origin='country.name', destination='country.name')
```

```
## Warning in countrycode(sourcevar = polity$country, origin = "country.name", : Some values were not matched
```

```
## Warning in countrycode(sourcevar = polity$country, origin = "country.name", : Some strings were matched
```

```
worldBank$cname <- countrycode(
  worldBank$country,
  'country.name', 'country.name')
```

```
## Warning in countrycode(worldBank$country, "country.name", "country.name"): Some values were not matched
```

```
head(unique(polity[,c('country','cname')]))
```

```
##      country      cname
## 161 Afghanistan Afghanistan
## 264   Albania   Albania
## 321   Algeria   Algeria
## 376   Angola   Angola
## 553  Argentina  Argentina
## 610   Armenia   Armenia
```

```
polity$cyyear <- paste(
  polity$cname, polity$year,
  sep='_'
)
worldBank$cyyear <- paste(
  worldBank$cname, worldBank$year,
  sep='_'
)
```

This matches the GDP growth variable, as we did in class

```
polity$gdpGrowth <- worldBank$gdpGrowth[
  match(
    polity$cyyear, worldBank$cyyear
```

```

    )
  ]

  matched <- head(match(polity$year, worldBank$year))

  polity$year[1:6]

  ## [1] "Afghanistan_1960" "Afghanistan_1961" "Afghanistan_1962"
  ## [4] "Afghanistan_1963" "Afghanistan_1964" "Afghanistan_1965"
  worldBank[matched, 'year']

  ## [1] "Afghanistan_1960" "Afghanistan_1961" "Afghanistan_1962"
  ## [4] "Afghanistan_1963" "Afghanistan_1964" "Afghanistan_1965"

  And finally, we use a loop to match the remaining variables
  worldBank <- worldBank[!is.na(worldBank$country),] # remove NAs (from additional regions that were in w

  to_match <- list("gdpConstant", "popGrowth", "tradeOpen")

  for(i in to_match){

    polity[,i] <- worldBank[,i] [
      match(
        polity$year, worldBank$year
      )
    ]
  }

  head(polity)

```

```

##      ccode scode      country year democ autoc polity2 xconst      cname
## 161    700   AFG Afghanistan 1960     0    10    -10      1 Afghanistan
## 162    700   AFG Afghanistan 1961     0    10    -10      1 Afghanistan
## 163    700   AFG Afghanistan 1962     0    10    -10      1 Afghanistan
## 164    700   AFG Afghanistan 1963     0    10    -10      1 Afghanistan
## 165    700   AFG Afghanistan 1964     0     7     -7      3 Afghanistan
## 166    700   AFG Afghanistan 1965     0     7     -7      3 Afghanistan
##
##      cyear gdpGrowth gdpConstant popGrowth tradeOpen
## 161 Afghanistan_1960      NA      NA  1.816077  11.15703
## 162 Afghanistan_1961      NA      NA  1.876528  12.55061
## 163 Afghanistan_1962      NA      NA  1.934999  14.22764
## 164 Afghanistan_1963      NA      NA  1.992521  26.03551
## 165 Afghanistan_1964      NA      NA  2.049423  26.94445
## 166 Afghanistan_1965      NA      NA  2.105369  32.67108

```

Problem 5

Hadley problems

Data structures:

Vectors: 1. What are the six types of atomic vector?

logical, integer, double, character, complex and raw How does a list differ from an atomic vector? lists can include any element type, including other lists

2. What makes `is.vector()` and `is.numeric()` fundamentally different to `is.list()` and `is.character()`? `is.vector` and `is.numeric` do not test for a specific type of element
3. Test your knowledge of vector coercion rules by predicting the output of the following uses of `c()`: `c(1, FALSE) <- double` `c("a", 1) <- character` `c(list(1), "a") <- list` `c(TRUE, 1L) <- integer`
4. Why do you need to use `unlist()` to convert a list to an atomic vector? Why doesn't `as.vector()` work? Lists are vectors, but not atomic vectors, so `as.vector` doesn't do anything as it is coercing a vector to a vector.
5. Why is `1 == "1"` true? Why is `-1 < FALSE` true? Why is `"one" < 2` false? `1 == "1"`: the numeric is coerced into a character - moving up from a less flexible element to more `-1 < F: FALSE` is coerced into an integer, become 0, making the statement `-1 < 0` `"one" < 2`: the numeric is coerced into a character and `"one" < "2"` is then FALSE
6. Why is the default missing value, NA, a logical vector? What's special about logical vectors? (Hint: think about `c(FALSE, NA_character_)`.) Logical is the least flexible type, so NA are logical so they don't lose any information when coerced.

Attributes: 1. An early draft used this code to illustrate `structure()`: `structure(1:5, comment = "my attribute")` [1] 1 2 3 4 5 But when you print that object you don't see the comment attribute. Why? Is the attribute missing, or is there something else special about it? (Hint: try using `help()`.) the "comment" attribute is a special kind that does not print, unlike the others

2. What happens to a factor when you modify its levels? `f1 <- factor(letters)` `levels(f1) <- rev(levels(f1))` It reversed both the levels and the factor
3. What does this code do? How do `f2` and `f3` differ from `f1`? `f2` reverses the factor, but not the levels `f3` reverses the levels, but not the factor

Matrices and Arrays 1. What does `dim()` return when applied to a vector? NULL

2. If `is.matrix(x)` is TRUE, what will `is.array(x)` return? TRUE, an array is a matrix
3. How would you describe the following three objects? What makes them different to `1:5`? `x1 <- array(1:5, c(1, 1, 5))` a 1x1x5 array `x2 <- array(1:5, c(1, 5, 1))` a 1x5x1 array `x3 <- array(1:5, c(5, 1, 1))` a 5x1x1 array They differ from `1:5` in that they have dimensions (3)

Data Frames: 1. What attributes does a data frame possess? `names`, `row.names`, `class`

2. What does `as.matrix()` do when applied to a data frame with columns of different types? Coerces everything into the same type.
3. Can you have a data frame with 0 rows? What about 0 columns? Yes, an empty data frame

Subsetting:

Data Types:

1. Fix each of the following common data frame subsetting errors: `mtcars[mtcars$cyl = 4,] -> mtcars[mtcars$cyl == 4,]` `mtcars[-1:4,] -> mtcars[-(1:4),]` `mtcars[mtcars$cyl <= 5] -> tcars[mtcars$cyl <= 5,]` `mtcars[mtcars$cyl == 4 | 6,] -> mtcars[mtcars$cyl == 4|mtcars$cyl == 6,]`
2. Why does `x <- 1:5; x[NA]` yield five missing values? (Hint: why is it different from `x[NA_real_]`?) `X` takes `NA` and expands it to a length of 5
3. What does `upper.tri()` return? How does subsetting a matrix with it work? Do we need any additional subsetting rules to describe its behaviour? `x <- outer(1:5, 1:5, FUN = "**")` `x[upper.tri(x)]` `upper.tri()` returns a logical matrix the size of the upper triangle of the matrix. Using it to subset the matrix, is in `x[upper.tri(x)]`, returns a vector of the subsetted part of the matrix
4. Why does `mtcars[1:20]` return an error? How does it differ from the similar `mtcars[1:20,]`? `mtcars[1:20]` tells R to look at 20 columns, but there are not that many columns in `mtcars`. `mtcars[1:20,]` tells R to look at rows 1:20 and all the columns
5. Implement your own function that extracts the diagonal entries from a matrix (it should behave like `diag(x)` where `x` is a matrix).

```
my_diag <- function(x){ n <- nrow(x) return(x[matrix(seq_len(n), nrow = n, ncol = 2)]) }
```

6. What does `df[is.na(df)] <- 0` do? How does it work? subsets a dataframe to the values that are `NA` and turns `NAs` into `0s`.

####Subsetting Operators

1. Given a linear model, e.g., `mod <- lm(mpg ~ wt, data = mtcars)`, extract the residual degrees of freedom. Extract the `R squared` from the model summary (`summary(mod)`) `mod$df.residual` `summary(mod)$r.squared`

####Applications

1. How would you randomly permute the columns of a data frame? (This is an important technique in random forests.) Can you simultaneously permute the rows and columns in one step? `df[,sample(ncol(df))]` `df[sample(nrow(df)),sample(ncol(df))]`
2. How would you select a random sample of `m` rows from a data frame? What if the sample had to be contiguous (i.e., with an initial row, a final row, and every row in between)? `m <- some number` `testdf[sample(nrow(testdf), size = m),]` `testdf[sample(nrow(testdf) - m + 1, size = 1) + (0:(m - 1)),]`
3. How could you put the columns in a data frame in alphabetical order? `testdf[order(names(testdf))]`